

Evaluation & Awards Ecosystem

Proposed by: [Name, Rank] · 96th Maintenance Squadron, 96 MXG · [Email / DSN]

Bottom line: Evaluations and awards are valuable. The hours spent writing, emailing, marking up and re-sending them are not. We propose a secure, AI-assisted ecosystem that **gives that administrative time back to the warfighter** and produces stronger records, piloted in the 96 MXS and built to scale across the DAF at IL4/IL5.

THE PROBLEM

Every quarter and every year, Airmen and Guardians rebuild their accomplishments from memory. Packages then move as PDFs by email: opened, marked up, saved and sent back at every layer. An estimated **60–80% of this work happens right at the deadline**. The know-how for a strong, promotable record lives in scattered PDFs and in a few experienced supervisors' heads. The result is lost detail, uneven quality, version confusion, and hours pulled away from the mission.

THE SOLUTION: THREE LAYERS, ONE AI ASSISTANT

- **Capture.** A year-round journal. Members log jobs, efforts, dates, impact and metrics in about 30 seconds, as work happens.
- **Context.** Each unit's rules are codified: who can apply for which award, format (e.g., 5 statements, 6 lines max), routing and suspense dates. The SEL sets it up in a **5–10 minute guided interview**, inheriting from group and wing. A **document library** of writing guides, award SOPs and sanitized past winners teaches the AI the unit's playbook.
- **Coordination.** One live package routes member > supervisor > section > flight > SEL > board, with inline comments, tracked edits, returns and full history. No email or PDF loop.
- **AI assistant.** Coaches members ("how many aircraft? what did it save?"), drafts from journal entries, checks format live, scores the *writing* against the award rubric, suggests edits, and speeds supervisor review. Every suggestion cites its source.

RESPONSIBLE AI BY DESIGN

- **Humans decide.** The AI drafts and advises and never submits.
- **AI scores the writing, not the person.** Boards never see AI scores and score packages with names removed.
- **Traceable.** Every AI output is logged with model and prompt version and traced to the member's own entries.
- **Private by default.** A member's journal is theirs. Supervisors see only submitted packages. Access follows the rating chain.

BUILT FOR DOD DEPLOYMENT

- Cloud-agnostic: **AWS GovCloud or Azure Government**, targeting **IL4** (CUI/PII) with an IL5 path.
- Kubernetes and Platform One compatible; Iron Bank images; FIPS crypto; CAC/PKI login; air-gap ready.
- Designed to inherit platform controls; Privacy Impact Assessment planned early.
- A prep and coordination tool that feeds official systems of record rather than replacing them.

ROADMAP

When	Phase	Exit gate
FY27 Q1	0 · Pitch & prototype. Clickable demo on synthetic 96 MXS data.	Selected; sponsor named
FY27 Q2	1 · MVP. Journal, library & import, SEL setup, rules & six-line check, routing, AI coach & editor.	Full cycle on synthetic data
FY27 Q3	2 · 96 MXS pilot. IL4 environment, privacy & ATO path, one real quarterly award cycle, AI scorer & copilot.	Pilot metrics met
FY27 Q4–FY28 Q1	3 · Group & wing. Nominations advance to 96 MXG and wing; dashboards; evaluation drafts.	New squadron onboarded in < 1 day
FY28+	4 · Enterprise. Other wings, both clouds, IL5 option, export to official systems.	Enterprise sponsor & sustainment

HOW WE'LL MEASURE SUCCESS

- **Admin hours returned to the mission** per package (member + every reviewer), measured against a pre-pilot baseline.
- Share of packages ready **before the final week**, versus today's deadline crunch.
- Emails and PDF versions eliminated; review rounds per package.
- Packages rule-compliant on first submission; member and supervisor satisfaction.

HOW THE AI FACTORY CAN HELP

1. **Sponsorship.** Accept the project into the AI Factory portfolio and help align 96 TW / 96 MXG leadership for the pilot.
2. **Environment.** An IL4 development and pilot environment on the AI Factory platform (AWS GovCloud or Azure Government), with a DevSecOps pipeline.
3. **AI model access.** An approved LLM and embedding endpoint authorized for CUI/PII workloads.
4. **Accreditation & privacy.** Guidance on control inheritance and the ATO path, Privacy Impact Assessment and SORN determination, and Responsible AI review.
5. **Team.** Engineering and UX support (proposed: 2–3 developers and 1 designer for two quarters) alongside the unit's domain experts.
6. **Resourcing.** Pilot funding and compute (ROM to follow after scoping with the AI Factory).

What the 96 MXS brings: a ready pilot unit with defined award programs and routing, SEL and supervisor subject-matter experts, the unit's writing guides and SOPs, and a working architecture and prototype.